

Self-RAG is a retrieval-augmented generation framework that trains a language model to retrieve passages on demand and generate reflection tokens. Reflection tokens let the model decide whether retrieval is needed and critique relevance, factual support, and response quality. The method integrates retrieval and self-reflection within generation.

Self-RAG controls retrieval adaptively during generation. Its reflection tokens provide explicit signals for retrieval necessity, passage relevance, claim support, and output utility. This design aims to improve evidence use without retrieving for every prompt.